

**Lab of EMIoT
Welcome!**

Objective and organization

- Logistics: In-class lab
 - 1 laptop per person
 - May be useful to bring portable multiple sockets
- Necessary software:
 - C
 - MATLAB
 - You can get a free student licence
https://www.areait.polito.it/supporto/risultato_serv.asp?serv=matlab&dettaglio=S&id_progetto_servizio=331

LAB schedules

- 20% of the final score
 - 9 points maximum
- Assignments will be evaluated
 - Each student has to deliver his/her own assignment
 - No groups allowed!
 - 1 report per lab
 - Any extension to the minimum assignment **may** lead to an increase in the evaluation
 - Make sure you meet all requirements
 - Do not go out of topic



LAB delivery

- Lab deadline is 23:59 of the day before the 2nd exam
 - **No exception**
 - Late delivery implies zero score for labs
- Format: **one archive**
 - File name = report.zip
 - One subfolder per lab (Lab1/ Lab2/ Lab3/)



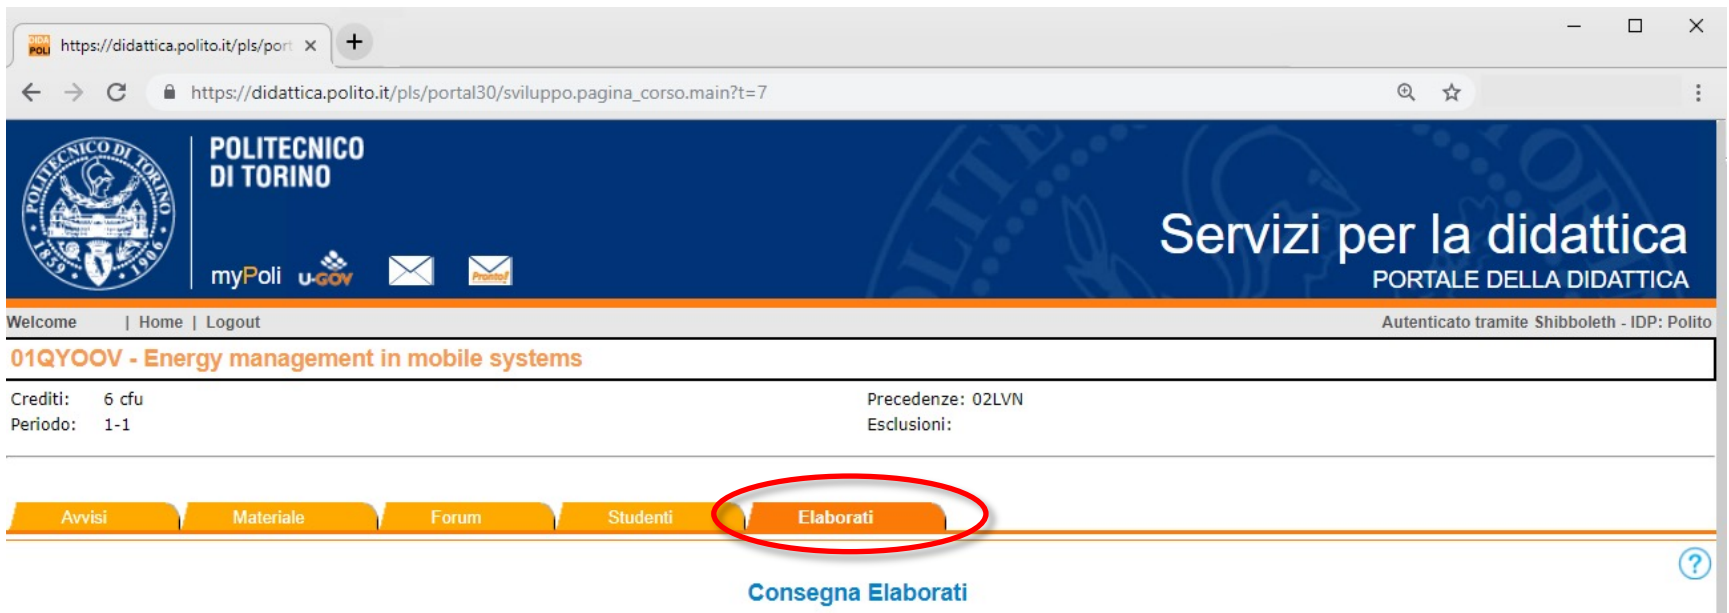
LAB delivery

- What to deliver for each lab:
 - Source code
 - All code you modified and/or consider necessary
 - Report
 - 5-10 pages per lab, depending on the depth of experiments
 - PDF format
 - **Analysis of results**
 - This is what gives you points!
 - Implementing the code is not enough!



LAB delivery

- How to deliver:
 - Through the didattica web site
 - «Elaborati» tag



The screenshot shows a web browser window displaying the Politecnico di Torino Didattica Portal. The browser's address bar shows the URL https://didattica.polito.it/pls/portal30/sviluppo.pagina_corso.main?t=7. The page header features the Politecnico di Torino logo and the text "Servizi per la didattica" and "PORTALE DELLA DIDATTICA". Below the header, there is a navigation bar with links: "Avvisi", "Materiale", "Forum", "Studenti", and "Elaborati". The "Elaborati" link is highlighted with a red circle. Below the navigation bar, the course information is displayed: "01QYOOV - Energy management in mobile systems". The course details include "Crediti: 6 cfu", "Periodo: 1-1", "Precedenze: 02LVN", and "Esclusioni:". At the bottom of the page, there is a link labeled "Consegna Elaborati".

https://didattica.polito.it/pls/portal30/sviluppo.pagina_corso.main?t=7

POLITECNICO DI TORINO

myPoli u-GOV

Servizi per la didattica
PORTALE DELLA DIDATTICA

Welcome | Home | Logout Autenticato tramite Shibboleth - IDP: Polito

01QYOOV - Energy management in mobile systems

Crediti: 6 cfu
Periodo: 1-1

Precedenze: 02LVN
Esclusioni:

Avvisi Materiale Forum Studenti **Elaborati**

Consegna Elaborati

Lab 1

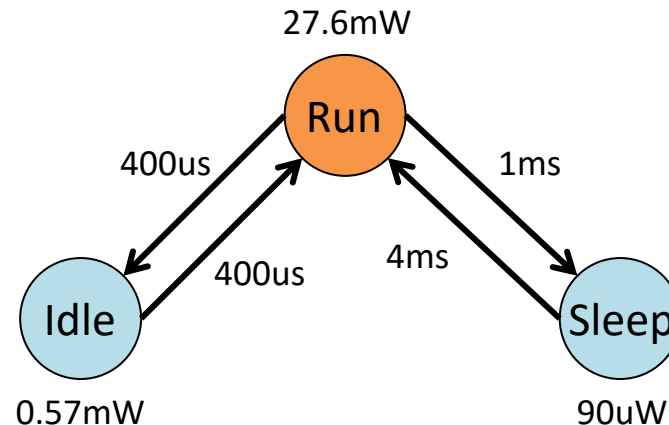
Dynamic Power Management

Objective and organization

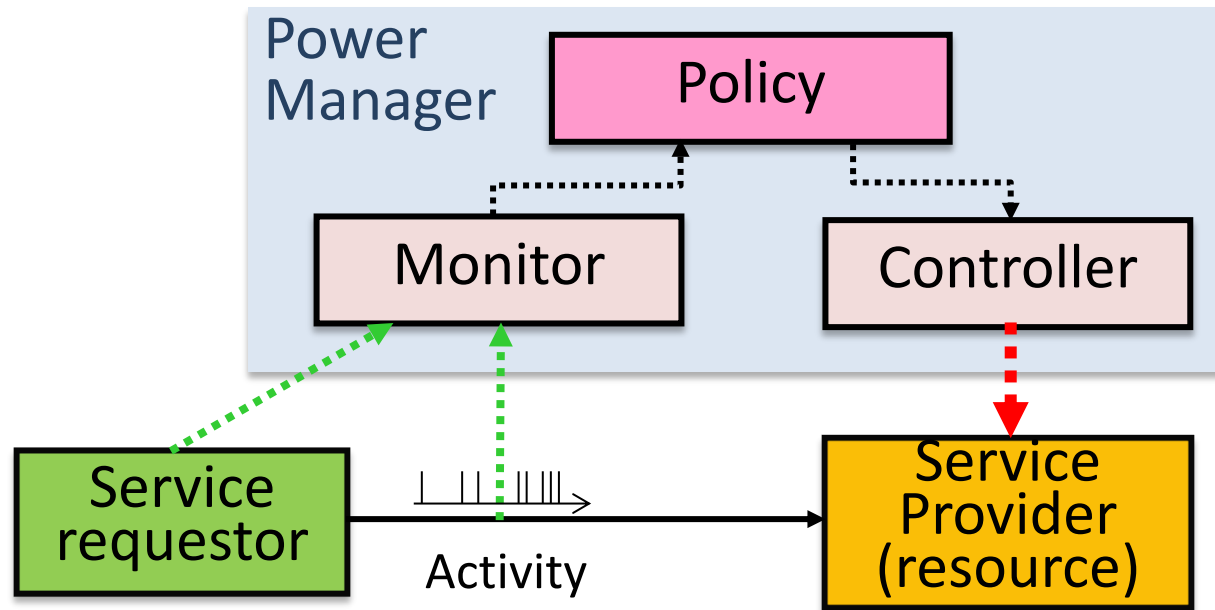
- Understanding of the basics of DPM
 - Use and modify a simple power state machine simulator in C
 - Evaluation of power management policies
- 1 report and 3 days

Recall

- Dynamic Power Management
 - Reduce power by turning devices to low power when peak performance is not needed
 - Devices abstracted as power state machines
 - Several internal states corresponding to modes of operation
 - Different power and service levels



Recall



- Power manager (PM)
 - Monitors requestor's activity and sets state of provider according to some **policy**
 - E.g., shuts down component after some inactivity time

Recall

- Given a PSM and a workload, determine the optimal allocation of power states over time that minimizes power under performance constraints
- Non-idealities make the problem non-trivial!
 - Transitions costs (time & power) are not zero
 - Transitions must be amortized!
 - Length of idle periods often unknown

DPM simulator

- Goal of the lab:
 - Evaluate on a case study how energy saving changes as a function of
 - The applied DPM policies
 - The distribution of idle times
 - The PSM parameters
 - ...

DPM simulator

- C program with the following basic operations
 - Read a power model → a PSM
 - Read a workload profile
 - Simulate two power management policies
 - Timeout
 - History-based prediction

DPM simulator

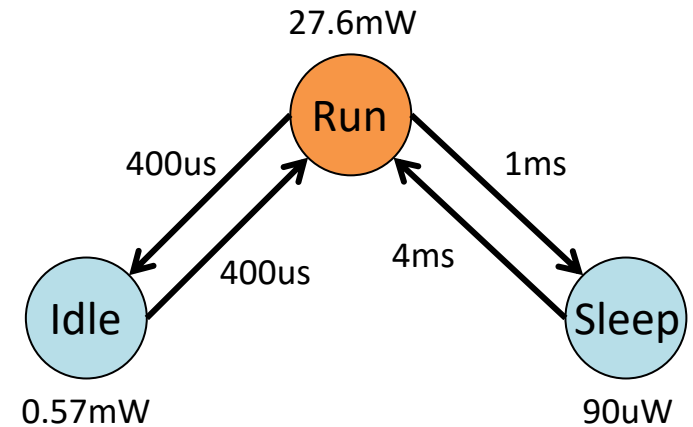
- `dpm_simulator [-help] [-t|-h] [-psm <psm file>] [-wl <wl file>]`
 - `-help`: prints command line instructions and returns
 - `-t <Timeout>`: use a *timeout* policy with <Timeout>
 - `-h <Value1> ...<ValueN>`: use a *history-based predictive policy*. <Value1-N> are additional policy parameters
 - `-psm <psm file>`: the name of the file describing the *power state machine* (PSM) of the resource
 - `-wl <wl file>`: the *workload* file name

Format of the PSM

27.6	0.57	0.09
0/0	0.01/0.4	0.02/1
0.01/0.4	0/0	-1/-1
2/4	-1/-1	0/0

States power

Transitions costs
(energy/time)



Transition does not exist:

- 0/0: Self-loops (i.e., state does not transition to itself)
- -1/-1: There is no transition between states

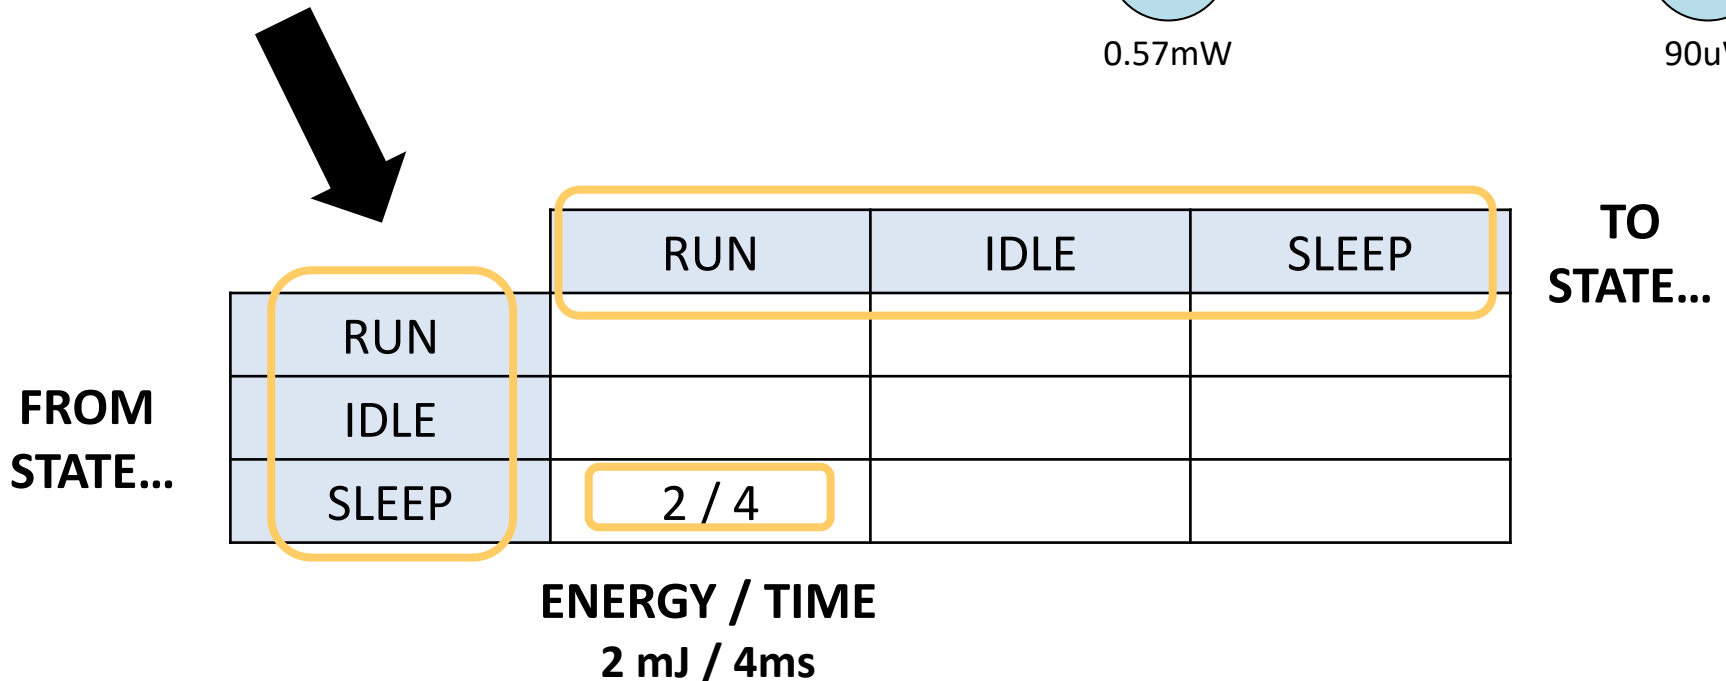
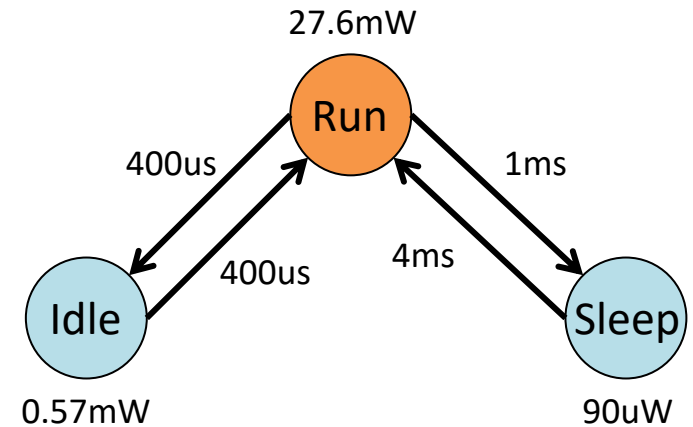
Default time, power, and energy units are: ms, mW and mJ

Format of the PSM

27.6	0.57	0.09
0/0	0.01/0.4	0.02/1
0.01/0.4	0/0	-1/-1
2/4	-1/-1	0/0

States power


Transitions costs
(energy/time)



Workload format

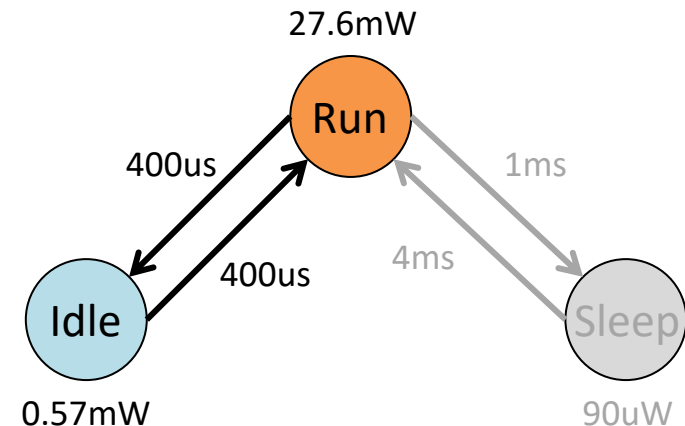

- The workload is given as a list of idle intervals
 - Values are in μs

268	341	START/END OF 1st IDLE PERIOD
609	656	START/END OF 2nd IDLE PERIOD
1010	1065	...
1196	1273	
1541	1629	
2056	2139	



DPM policies

- Timeout policy
 - Observe the first part of the current idle period to predict the length of the remaining part
 - Put the device in off state T_{TO} time units after it has entered the idle state



Compile and execute

- Compile (requires gcc):

`make`

- Generate Documentation (requires doxygen):

`doxygen Doxyfile`

- Generates «docs» folder with HTML documentation

- Execute:

`./dpm_simulator -t 20 -psm example/psm.txt -wl
example/wl.txt`

- Timeout policy with timeout value 20ms

Compile and execute


- Compile (requires gcc):

```
(base) → dpm-simulator git:(master) x ./dpm_simulator -psm example/psm.txt -wl example/wl.txt -t 20
[psm] State Run: power = 27.6000mW
[psm] State Idle: power = 0.5700mW
[psm] State Sleep: power = 0.0900mW
[psm] Run -> Idle transition: energy = 0.0100mJ, time = 0.4000ms
[psm] Run -> Sleep transition: energy = 0.0200mJ, time = 1.0000ms
[psm] Idle -> Run transition: energy = 0.0100mJ, time = 0.4000ms
[psm] Sleep -> Run transition: energy = 2.0000mJ, time = 4.0000ms

[sim] Active time in profile = 300.130000s
[sim] Inactive time in profile = 244.066000s
[sim] Total time = 544.196000s
[sim] Timeout waiting time = 24.679000s
[sim] Total time in state Run = 324.809000s
[sim] Total time in state Idle = 219.387000s
[sim] Total time in state Sleep = 0.000000s
[sim] Time overhead for transition = 0.910800s
[sim] N. of transitions = 2277
[sim] Energy for transitions = 0.0227700000J
[sim] Energy w/o DPM = 15.0198095999J, Energy w DPM = 9.1125489900J
(base) → dpm-simulator git:(master) x
```

Workloads

1. Two workloads provided on the course page
 - Similar structure (typical IoT workloads)
 - Read some data from sensors (polling)
 - Elaborate the data (e.g. Neural Network inference)
 - Repeat every 2 minutes

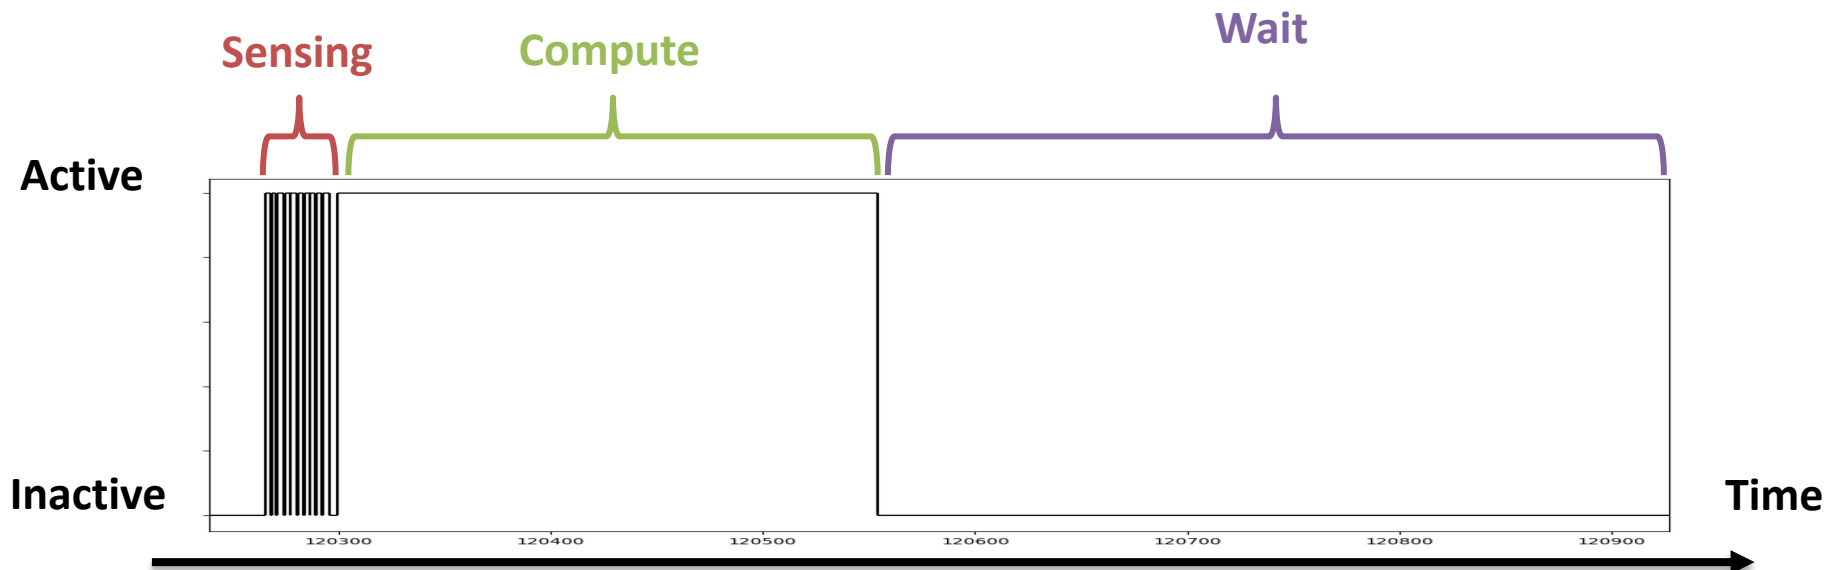


Workloads

1. Two workloads provided on the course page
 - Workloads generated using an embedded system simulator called JUMPER: <https://docs.jumper.io/>
 - You **don't have** to use JUMPER directly for the lab, but you can (if you want) use it to generate additional workloads
 - Instructions and a basic simulation setup are provided on the Polito website.

Workloads

1. Two workloads provided on the course page
 - workload_1.txt: "fast" sensors → 4ms to return a value
 - workload_2.txt: "slow" sensors → 100ms to return a value
 - **Does this have an impact on DPM?** It's your job to find it out




Part 1

Default Timeout Policy

Assignment 1 – Part 1

1. Compile the DPM simulator
2. Test it with the two workloads, using the default timeout policy
 - Only transitions between RUN and IDLE states
 - Try different timeout values and see what happens. Discuss in your report why it happens!



Part 2

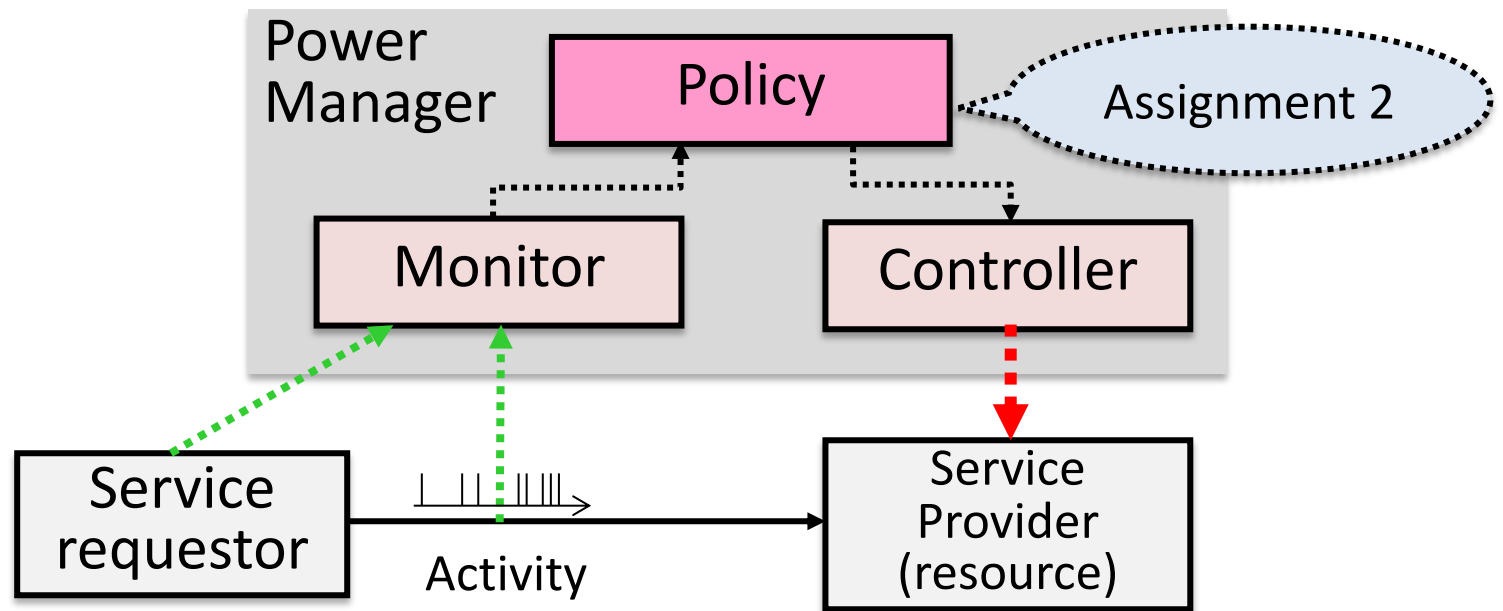
Extension of the timeout policy

DPM simulator

- The DPM simulator supports a timeout policy with transitions from ACTIVE to IDLE
 - Never goes to SLEEP state
 - Moving to SLEEP may save even more energy...

```
[sim] Active time in profile = 300.130000s
[sim] Inactive time in profile = 244.066000s
[sim] Total time = 544.196000s
[sim] Timeout waiting time = 24.679000s
[sim] Total time in state Run = 324.809000s
[sim] Total time in state Idle = 219.387000s
[sim] Total time in state Sleep = 0.000000s
[sim] Time overhead for transition = 0.910800s
[sim] N. of transitions = 2277
[sim] Energy for transitions = 0.0227700000J
[sim] Energy w/o DPM = 15.0198095999J, Energy w
DPM = 9.1125489900J
```



Recall



- Policy implementation
 - Modify the timeout-based policy

Assignment 1 – Part 2

- Modify the timeout policy to enable transitions also to SLEEP
 - Must modify the implementation of the DPM simulator



```
switch (policy) {
```

```
case DPM_TIMEOUT:
```

```
/* Day 2: EDIT */
```

```
if(curr_time > idle_period.start+tparams.timeout) {
```

```
    *next_state = PSM_STATE_IDLE;
```

```
} else {
```

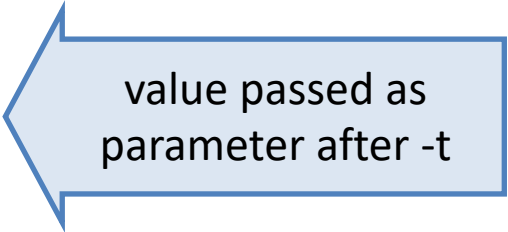
```
    *next_state = PSM_STATE_ACTIVE;
```

```
}
```

```
break;
```

The DPM simulator

- `print_command_line()` (src/utilities.c)
 - Prints the command line to invoke the tool
- `parse_arg()` (src/utilities.c)
 - Parses the inputs you provide via command line
 - For the timeout policy:
 - *selected_policy = DPM_TIMEOUT;
 - tparams->timeout = atof(argv[++cur]);
 - May have to be modified, too



value passed as
parameter after -t

The DPM simulator

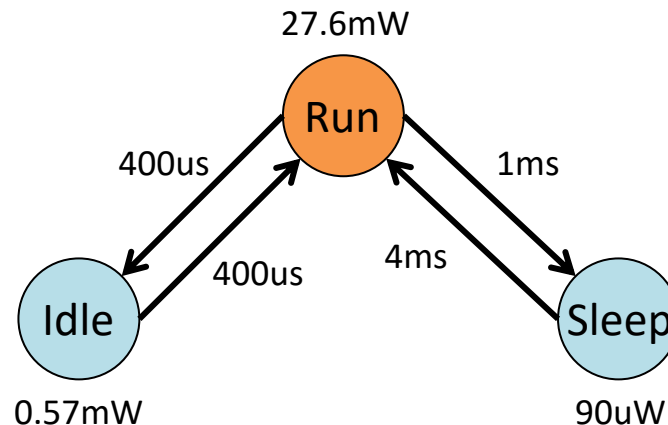
- `dpm_decide_state()` (src/dpm_policies.c)
 - Provided system state
 - Workload status (active/idle)
 - Current simulation time
 - Adopted policy (i.e., timeout)
 - Determines what is the state at the next time step

The DPM simulator

- `dpm_simulate()` (src/dpm_policies.c)
 - Emulates policy application to the PSM given the workload
 - For all instants in the current (active and idle) period
 - Invoke `dpm_decide_state()` to apply the policy and determine transitions
 - Check the state returned and update energy consumption accordingly


Power State Machine

- Note that there is **no transition** between IDLE and SLEEP in the PSM
 - This system can only go to IDLE **or** SLEEP **from the RUN state.**
 - With a timeout policy **only one of the two low-power states can be selected**
 - The story is different for predictive policies.

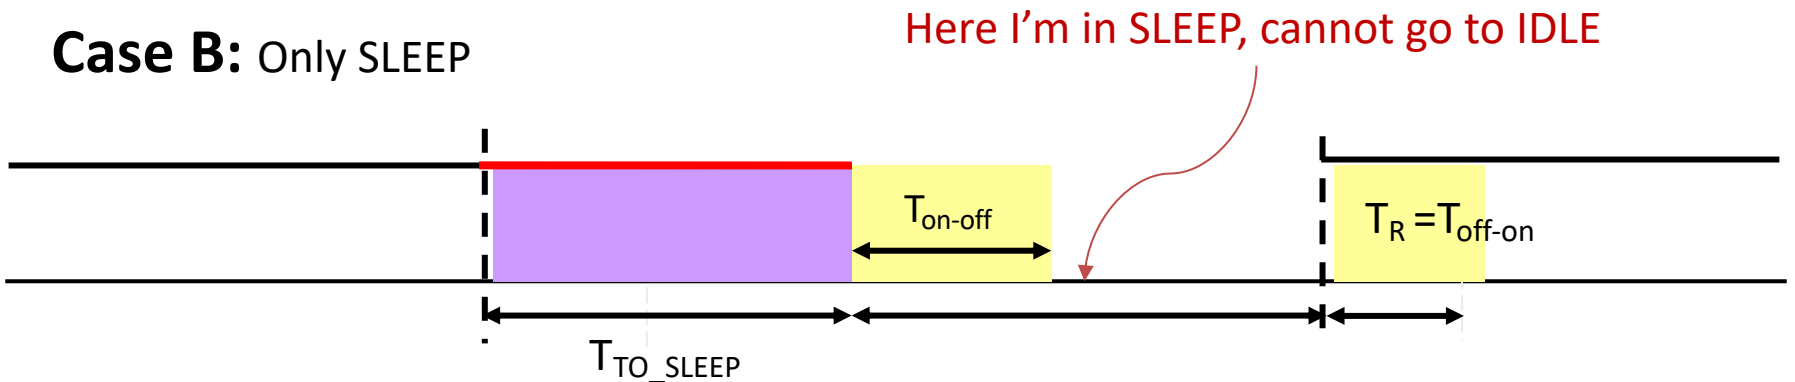


Timeout Policy

Case A: Only IDLE




Case B: Only SLEEP



NOTE: There is actually a way to use both low-power states.
Which one? [Possible Extra]

Assignment 1 – Part 2


- Report assignment:
 - Comparison between RUN-> IDLE timeout policy and RUN->SLEEP timeout policy
 - What changes?
 - What's the **best** T_{TO} value in the two cases?
 - Which of the two results in the overall lower power? What changes for the two **workloads** provided?
 - What changes in terms of the **timing overheads**?
 - **Why?**



This is the most important!!!

Assignment 1 – Part 2

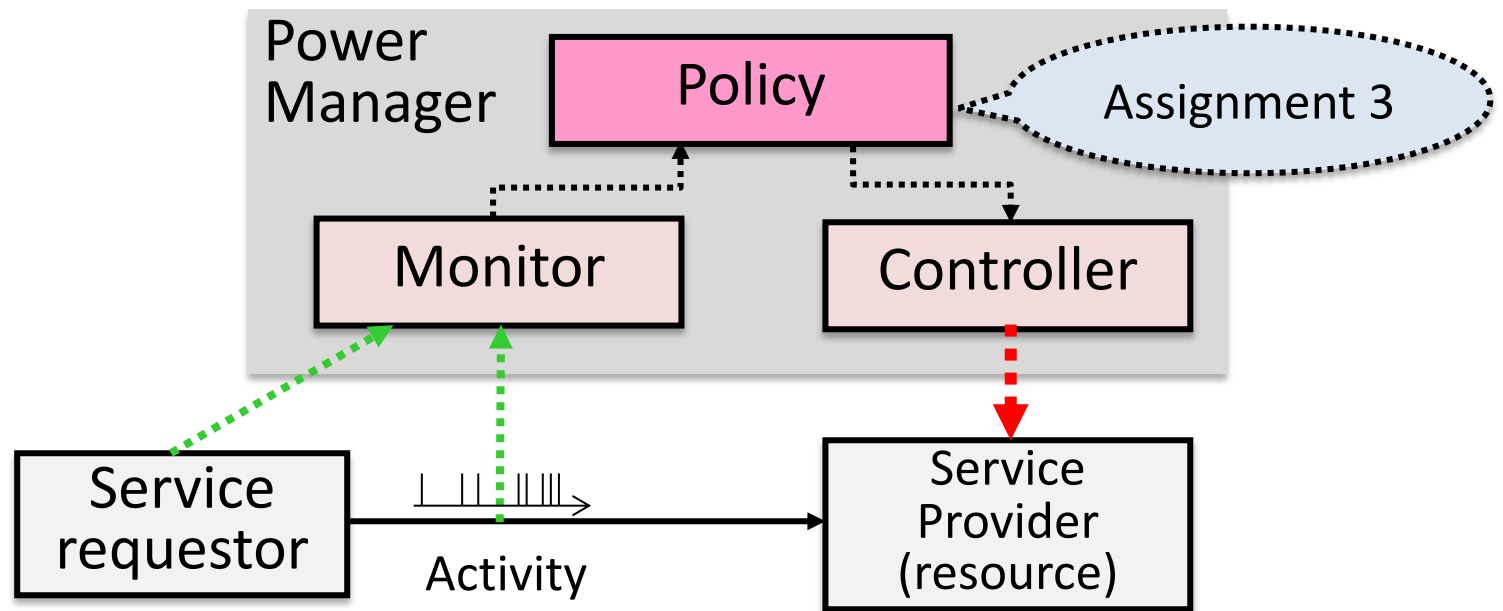
- **Extra: Make your analysis automatic and systematic**
 - Don't just try some “random” T_{TO} values
 - Compare things in a reasonable and meaningful way



Part 3

Predictive Policy

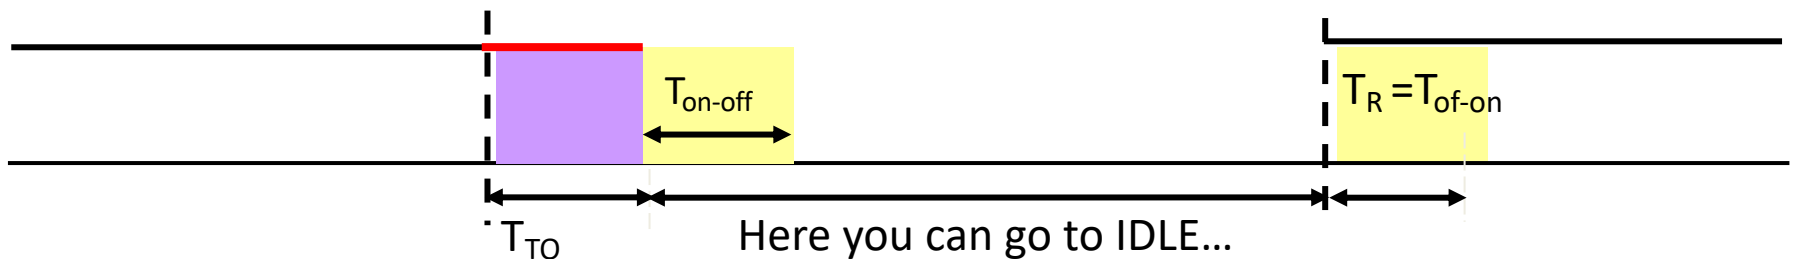
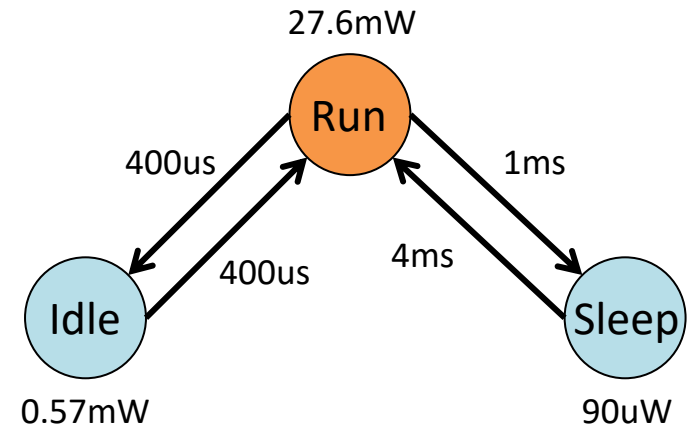
Recall



- Policy implementation
 - Implement a predictive policy

Recall

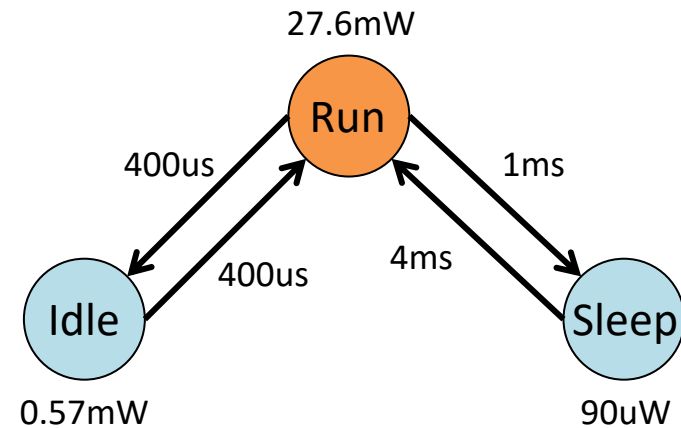
- So far we worked with timeout policies...
 - Put the device in off state T_{TO} time units after it has entered the idle state



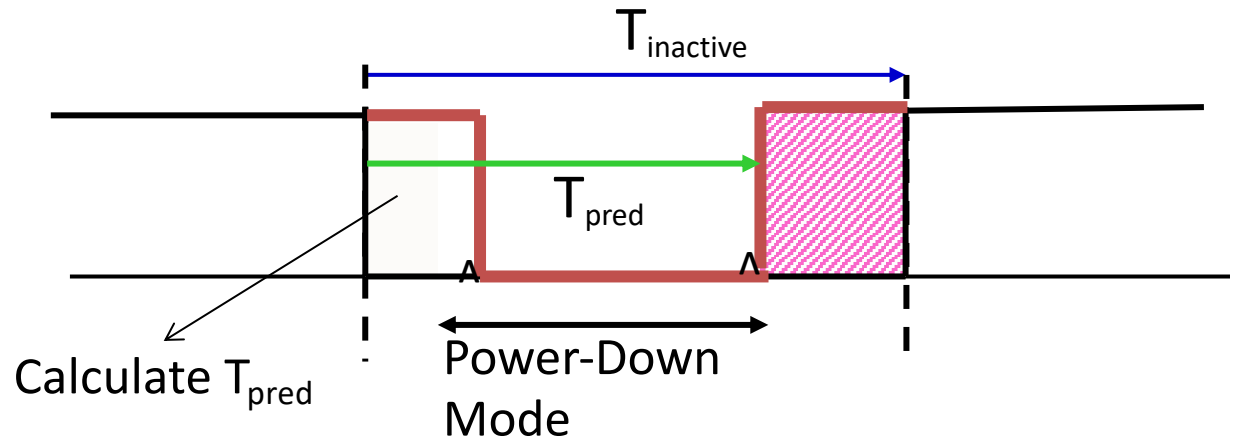
Recall

- So far, we worked with timeout policies...

- Put the device in off state T_{TO} time units after it has entered the idle state



- Can we do better?



Recall


- Predictive policies
 - Predict inactive period $T_{\text{pred}} \sim T_{\text{idle}}$
 - Go to sleep state if T_{pred} is long enough to amortize state transition cost
- History-based policies:
 - Predictive policies that use previous **history** of T_{active} **and/or** T_{idle}
 - Example: non-linear regression equation



$$T_{\text{idle}}[i] = K + K_1 \cdot T_{\text{idle}}[i-1] + K_2 \cdot T_{\text{active}}[i] + K_3 \cdot T_{\text{active}}[i]^2$$

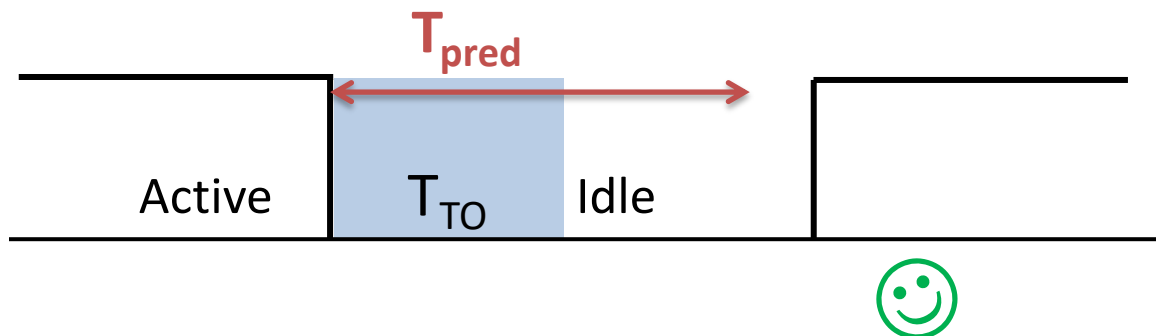
Recall

- Goal of predictive policies



Recall

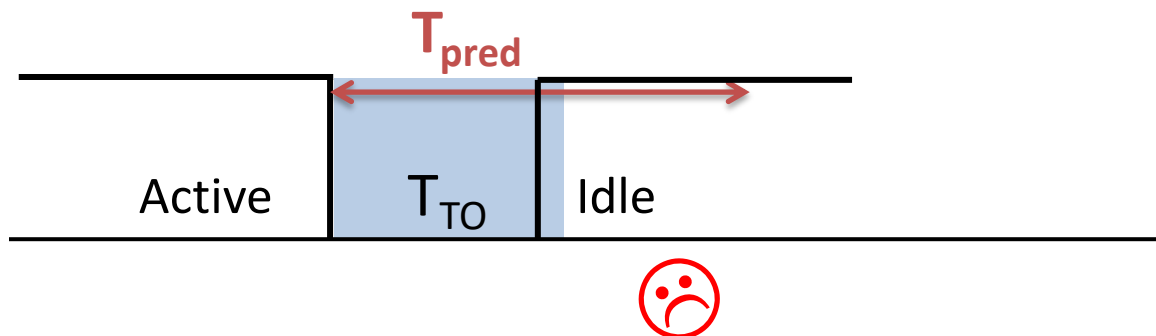
- Drawback of predictive policies
 - Under-prediction



Drawback: will my guess be right?

Recall


- Drawback of predictive policies
 - Over-prediction



Drawback: will my guess be right?

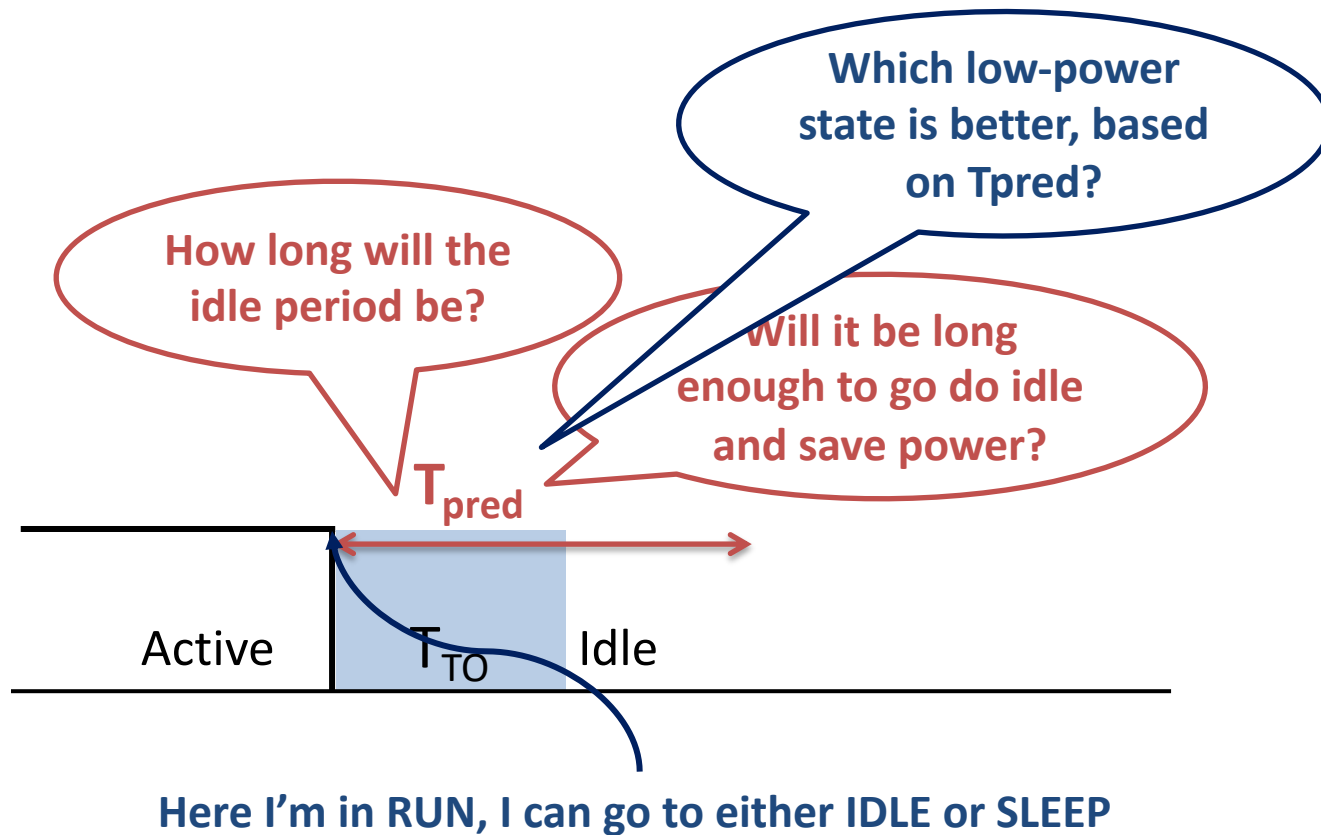
Recall

- But remember our workloads:
 - Pretty easy to predict (roughly) the duration of the next inactive phase.
 - How?



Recall

- A predictive policy also allows us to use both low-power states in the PSM:



The DPM Simulator

- Currently the simulator contains data structures and functions to implement a **regression-based** policy:

- Something like:

$$T_{idle}[i] = K + K_1 \cdot T_{idle}[i-1] + K_2 \cdot T_{active}[i] + K_3 \cdot T_{active}[i]^2$$

- *Mainly for «historical reasons»*
- **May not be the best type** of history-based policy to implement

The DPM Simulator

```
int parse_args(...) {
```

```
    //...
```

```
    if(strcmp(argv[cur], "-h") == 0) {
```

```
        *selected_policy = DPM_HISTORY;
```

```
        if(argc > cur + DPM_HIST_WIND_SIZE + 2){
```

```
            int i;
```

```
            for(i = 0; i < DPM_HIST_WIND_SIZE; i++) {
```

```
                hparams->alpha[i] = atof(argv[++cur]);
```

```
            }
```

```
            hparams->threshold[0] = atof(argv[++cur]);
```

```
            hparams->threshold[1] = atof(argv[++cur]);
```

```
        } else return 0;
```


```
    }
```


The DPM Simulator

```
/* update idle time history */  
void dpm_update_history(...) {  
    for (int i=0; i<DPM_HIST_WIND_SIZE-1; i++){  
        h[i] = h[i+1];  
    }  
    h[DPM_HIST_WIND_SIZE-1] = new_idle;  
}
```

The DPM Simulator

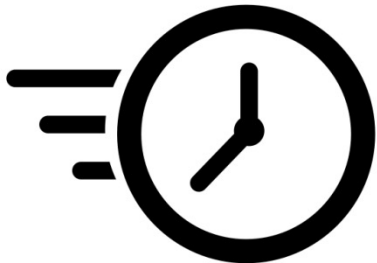
- Do we really need a polynomial?
- ...or can we make our decision simpler?
 - What determines if the next inactive period will be long or short?



Assignment 1 – Part 3

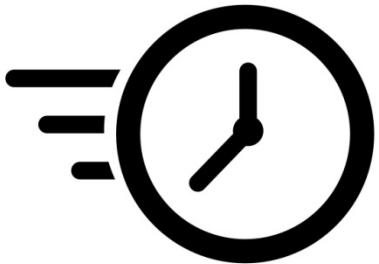
- Modify the simulator to implement a *predictive policy*
 - It can be **of any kind**, you decide
 - Trying and comparing more than one policy is also good
 - Of course, motivate your choices.

```
case DPM_HISTORY:  
    /* Day 3: EDIT */  
    if(curr_time < idle_period.start) {  
        *next_state = PSM_STATE_ACTIVE;  
    } else {  
        *next_state = PSM_STATE_ACTIVE;  
    }  
    break;
```



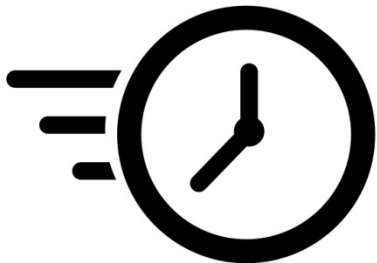
Assignment 1 – Part 3

- Report assignment
 - Description of implemented predictive policy
 - Result of implemented predictive policy with the workload profiles
 - Analysis on effect of policy parameters (if any)



Assignment 1 – Part 3

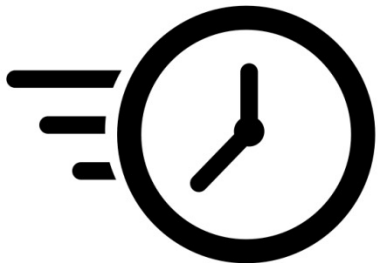
- Report assignment:
 - Comparison between predictive policy and timeout policies
 - What changes?
 - Which approach is the best for the two workloads?
 - Why does a predictive approach work better/worse on one workload than the other (if it does)?
 - What changes in terms of the **timing overheads**?
 - **Why?**



This is the most important!!!

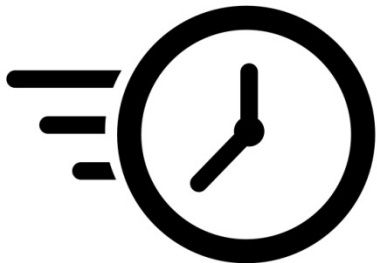
Assignment 1 – Part 2

- **Extras:**
 - Any extra analysis/experiment is well appreciated
 - Show your creativity and desire to explore the topic
 - Extras are **not** mandatory, but can give additional points:
 - If you only do the mandatory points of the previous slides, perfectly, you'll still get max score



Assignment 1 – Part 2

- **Extras: some suggestions**
 - Change the PSM (i.e., the target system)
 - Change the workload:
 - Generate new “synthetic” workloads using any tool of your choice (MATLAB, Python, etc).
 - You can simply generate the sequence of inactive intervals, but they should make sense.
 - Install and use **Jumper** to generate workloads starting from real embedded code.
 - Try other policies...
 - Etc.



End of Lab 1!
**Now you're ready to prepare
the first report...**

The deadline is 23:59 of the day before the 2nd exam